

Likelihood Optimization to Linear Regression with R

Fernando Delgado

Contents

| | |
|-----------------------------|---|
| Introduction | 1 |
| Loading the Data | 1 |
| Linear Regression | 2 |
| Validation | 3 |
| References | 4 |

Introduction

Linear regression is a statistic model used for predicting a numerical quantity. The parameters of a linear regression model can be estimated using least squares or by maximum likelihood optimization. Maximum likelihood estimation is a probabilistic framework for automatically finding the probability distribution for the observed data.

Through this report, we perform a Likelihood Optimization to linear regression with R.

Loading the Data

First, we simulate a dataframe to work with:

```
# Create a first column to test
a <- rnorm(1000,0)
df <- a

# Create 5 columns with random data
columns <- cbind(1:5)

# For loop data into columns
for (i in columns){
  name <- paste("var",i, sep="")
  tmp <- rnorm(1000,0)
  df <- data.frame(df, tmp)
  names(df)[names(df) == 'tmp'] <- name
}
```

We create a dataframe with 1000 observations to 5 predictor variables using R's `rnorm()` assignin a mean of 0 to obtain normalized random data.

Then, to give a little bit of sense to our data, we use `USArrests` built-in data set and `fakeR`'s `simulate_data()` to obtain a random target variable of Arrests by Assaults by city. We obtain 1000 random observations that

we append to our previous dataset. The final result is a data frame with 1000 observations for 5 random predictor variables to 1 target variable (number of arrests by assault by city).

```
## Classes 'data.table' and 'data.frame': 1000 obs. of 6 variables:
## $ var1 : num 1.2507 -0.4434 -0.0904 -0.9828 0.2991 ...
## $ var2 : num 1.023 -0.76 -0.799 0.511 -0.803 ...
## $ var3 : num 0.465 2.447 -0.347 -0.721 1.491 ...
## $ var4 : num 0.686 -1.393 -0.548 -0.859 -0.607 ...
## $ var5 : num -0.00569 2.93659 -0.43356 0.32744 0.44956 ...
## $ Assault: num 187 171 242 160 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

Linear Regression

To optimize a function, it is important to consider what function we should optimize. To work with a linear regression function in this case, we know that it follows a normal distribution with a mean of 0 and an unknown standard deviation.

$$\sum_{i=1}^{i=n} R_i \tilde{N}(0, s)$$

Where R equals:

$$R_i = y_i - \hat{y}_i$$

Therefore, the objective is obtaining a function of y_i which minimizes the residuals and shows the best coefficients for our function.

With this in mind, we write an optimization function with the following code:

```
#Define likelihood function to optimise
ll_lm <- function(par, y, x1, x2, x3, x4, x5){

  alpha <- par[1]
  beta1 <- par[2]
  beta2 <- par[3]
  beta3 <- par[4]
  beta4 <- par[5]
  beta5 <- par[6]
  sigma <- par[7]

  R = y - alpha - beta1 * x1 - beta2 * x2 - beta3 * x3 - beta4 * x4 - beta5 * x5

  -sum(dnorm(R, mean = 0, sigma, log = TRUE))
}
```

Alpha refers to our target variable, and the 5 betas represent each predictor variable.

In order to use the `optim()` function, it must have `par` arguments. `Par` arguments need a vector with guesses for all unknown parameters. In our code above, `par` arguments include initial values in all 7 of the unknown parameters.

It is important to mention that by using `dnorm()` we obtain logarithmic values. This helps to sum the single likelihood values instead of the product.

The linear model we are fitting looks like this:

$$E(Y|X) = \alpha + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5$$

Therefore, the residuals are calculated like this:

$$R = y - \alpha - \beta_1 x_1 - \beta_2 x_2 - \beta_3 x_3 - \beta_4 x_4 - \beta_5 x_5$$

Moreover, since residuals are following a normal distribution with a mean of 0, what's left is to find the standard deviation that best fits our data. Hence, we minimize the sum of errors with `optim()` command by using a minus sign before the sum.

However, before running the optimization, we also estimate our coefficients by simply calculating the mean of each variable:

```
#Estimate Betas
est_alpha <- mean(fake_arrests$Assault)
est_beta1 <- mean(fake_arrests$var1)
est_beta2 <- mean(fake_arrests$var2)
est_beta3 <- mean(fake_arrests$var3)
est_beta4 <- mean(fake_arrests$var4)
est_beta5 <- mean(fake_arrests$var5)
est_sigma <- sd(fake_arrests$Assault)
```

To keep it simple, we do it manually, but this could be looped for a larger dataset.

Moving forward, we optimize our model searching our maximum likelihood estimates for the different coefficients. We introduce our function `ll_lm()`, the estimated coefficients and the data to be optimized:

```
mle_par <- optim(fn = ll_lm,
                par = c(alpha = est_alpha,
                        beta1 = est_beta1,
                        beta2 = est_beta2,
                        beta3 = est_beta3,
                        beta4 = est_beta4,
                        beta5 = est_beta5,
                        sigma = est_sigma),
                y = fake_arrests$Assault,
                x1 = fake_arrests$var1,
                x2 = fake_arrests$var2,
                x3 = fake_arrests$var3,
                x4 = fake_arrests$var4,
                x5 = fake_arrests$var5)
```

Finally, we obtain our estimated coefficients

```
##      alpha      beta1      beta2      beta3      beta4      beta5      sigma
## 153.209290 -1.310258   1.612228   1.133520  -1.654349   5.703489  74.355542
```

Validation

If we compare the estimate with the result of the `lm()` command for the same model, we observe some slight differences in the coefficients. However, since they are rather small it is probably due to our initial guesses for the parameters.

```
##
## Call:
## lm(formula = Assault ~ var1 + var2 + var3 + var4 + var5, data = fake_arrests)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -161.10  -48.62  -14.40   45.97  167.23
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  153.190      2.363   64.816 <2e-16 ***
## var1         -1.419      2.341   -0.606  0.5445
## var2          1.573      2.317    0.679  0.4974
## var3          1.029      2.363    0.436  0.6633
## var4         -1.766      2.364   -0.747  0.4553
## var5          5.695      2.327    2.448  0.0145 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 74.61 on 994 degrees of freedom
## Multiple R-squared:  0.007842,    Adjusted R-squared:  0.002851
## F-statistic: 1.571 on 5 and 994 DF,  p-value: 0.1654
```

References

Creating fake simulated data was inspired from this following post:

<https://rviews.rstudio.com/2020/09/09/fake-data-with-r/>

Maximization of likelihood was inspired by this following post:

https://www.joshua-entrop.com/post/optim_linear_reg/